

# Deep Neural Networks - Lecture 3

Marcin Mucha

March 15, 2017

The perceptron


Neural Networks and Backpropagation

Trouble with neural networks

# Where it all begins

- ▶ Many settings where humans outperform best algorithms.
- ▶ Idea: Emulate the human brain.
- ▶  $10^{11}$  neurons,  $10^{14}$  connections in the brain. Trouble?
- ▶ Compare: a 10 TB drive can store  $10^{14}$  bytes.


# Where it all begins




# The perceptron

The *perceptron* is a function that, given inputs  $x_1, \dots, x_k$  outputs  $[\sum_i w_i x_i + b \geq 0]$  (or a specific GD-like algo for fitting these).

Think:  $x_i$  are factors contributing to a decision, they are weighted and the perceptron makes a decision automatically!




# The perceptron



## Mark I Perceptron (1957)

- ▶ Funded by US Navy.
- ▶ Photo recognition.
- ▶ 20x20 photo-cells.
- ▶ Weights in potentiometers.
- ▶ Trained with the perceptron algorithm - weights updated by electric motors.

# The perceptron



## Mark I Perceptron (1957)

- ▶ Funded by US Navy.
- ▶ Photo recognition.
- ▶ 20x20 photo-cells.
- ▶ Weights in potentiometers.
- ▶ Trained with the perceptron algorithm - weights updated by electric motors.

*...the embryo of an electronic computer that [the Navy] expects will be able to walk, talk, see, write, reproduce itself and be conscious of its existence...*



# Training the perceptron

Impossible to train directly using gradient methods since gradient is 0 where it exists. Instead...



# Training the perceptron


Impossible to train directly using gradient methods since gradient is 0 where it exists. Instead...



Train using a smooth approximation (not exactly what the perceptron algorithm was doing!).

# Training the perceptron

Impossible to train directly using gradient methods since gradient is 0 where it exists. Instead...



Train using a smooth approximation (not exactly what the perceptron algorithm was doing!).

Many smooth step functions possible, e.g. the sigmoid function  $\sigma(x) = \frac{1}{1+e^{-x}}$ . Rings a bell?

# The power of perceptrons

Perceptrons can only separate half-spaces.

# The power of perceptrons

Perceptrons can only separate half-spaces.


**Exercise:** Show that for binary inputs, it can compute AND.

# The power of perceptrons

Perceptrons can only separate half-spaces.

**Exercise:** Show that for binary inputs, it can compute AND.

**Solution:**




# The power of perceptrons

Perceptrons can only separate half-spaces.

**Exercise:** Show that for binary inputs, it can compute AND.

**Solution:**



**Exercise:** Can a perceptron compute XOR?


# Network of perceptrons

**Exercise:** Show that a DAG of perceptrons can compute XOR.

# Network of perceptrons

**Exercise:** Show that a DAG of perceptrons can compute XOR.

**Solution:**






# Network of perceptrons

**Exercise:** Show that a DAG of perceptrons can compute XOR.

**Solution:**




**Exercise:** Show that a DAG of perceptrons can compute any boolean function.

# Network of perceptrons

**Exercise:** Show that a DAG of perceptrons can compute XOR.

**Solution:**



**Exercise:** Show that a DAG of perceptrons can compute any boolean function.

Note that this DAG can be very large (think: SAT).


# Training a DAG of perceptrons

...but can we train DAGs of perceptrons?

# Training a DAG of perceptrons

...but can we train DAGs of perceptrons?

More specifically: Given a DAG of sigmoid approximations to the perceptron, can we compute gradients of the loss function?



# Training a DAG of perceptrons


## Theorem

*We can compute the derivatives of the loss function over all node values in linear time.*

# Training a DAG of perceptrons

## Theorem


*We can compute the derivatives of the loss function over all node values in linear time.*



# Training a DAG of perceptrons

## Theorem


*We can compute the derivatives of the loss function over all node values in linear time.*



# Training a DAG of perceptrons

## Theorem

*We can compute the derivatives of the loss function over all node values in linear time.*






# Training a DAG of perceptrons

## Theorem


*We can compute the derivatives of the loss function over all node values in linear time.*



# Training a DAG of perceptrons

## Theorem


*We can compute the derivatives of the loss function over all node values in linear time.*



# Training a DAG of perceptrons

## Theorem

*We can compute the derivatives of the loss function over all node values in linear time.*



# Backpropagation


It seems that for a node  $z$ , we have

$$\frac{\partial L}{\partial z} = \sum_{z=z_0 \rightarrow \dots \rightarrow z_l=L} \prod_{i=1}^{l-1} \frac{\partial z_{i+1}}{\partial z_i}.$$

# Backpropagation

It seems that for a node  $z$ , we have


$$\frac{\partial L}{\partial z} = \sum_{z=z_0 \rightarrow \dots \rightarrow z_l=L} \prod_{i=1}^{l-1} \frac{\partial z_{i+1}}{\partial z_i}.$$



# Backpropagation

It seems that for a node  $z$ , we have


$$\frac{\partial L}{\partial z} = \sum_{z=z_0 \rightarrow \dots \rightarrow z_l=L} \prod_{i=1}^{l-1} \frac{\partial z_{i+1}}{\partial z_i}.$$



# Backpropagation

It seems that for a node  $z$ , we have

$$\frac{\partial L}{\partial z} = \sum_{z=z_0 \rightarrow \dots \rightarrow z_l=L} \prod_{i=1}^{l-1} \frac{\partial z_{i+1}}{\partial z_i}.$$




# Backpropagation

It seems that for a node  $z$ , we have

$$\frac{\partial L}{\partial z} = \sum_{z=z_0 \rightarrow \dots \rightarrow z_l=L} \prod_{i=1}^{l-1} \frac{\partial z_{i+1}}{\partial z_i}.$$

$$\frac{\partial v}{\partial s} \cdot \frac{\partial L}{\partial v} + \frac{\partial u}{\partial s} \cdot \frac{\partial v}{\partial u} \cdot \frac{\partial L}{\partial v}$$





# Backpropagation

By doing topological sorting and going from the back we can compute  $\frac{\partial L}{\partial z}$  for all computation nodes  $z$ .

# Backpropagation

By doing topological sorting and going from the back we can compute  $\frac{\partial L}{\partial z}$  for all computation nodes  $z$ .

But we want derivatives over parameters, not nodes!

# Backpropagation

By doing topological sorting and going from the back we can compute  $\frac{\partial L}{\partial z}$  for all computation nodes  $z$ .


But we want derivatives over parameters, not nodes!

Computing the derivatives of nodes over their parameters is trivial.


# Multilayer Perceptron (MLP)

- ▶ *Multilayer Perceptron (MLP)*, or a *Feedforward Neural Network* consists of several layers of units (sigmoid or other), each connected with previous and next.
- ▶ Number of layers (excluding inputs) is called the *depth*.
- ▶ The inner layers are called *hidden*.

# Multilayer Perceptron (MLP)



# Matrix algebra - notation



# Matrix algebra

Regular structure of MLPs facilitates algebraic shortcuts both in forward and backward pass.

# Matrix algebra

Regular structure of MLPs facilitates algebraic shortcuts both in forward and backward pass.

$$f^k = W^k g^{k-1}.$$



# Matrix algebra

Regular structure of MLPs facilitates algebraic shortcuts both in forward and backward pass.

$$f^k = W^k g^{k-1}.$$

Here  $f^k$  and  $g^{k-1}$  could be vectors if single datapoint, or matrices.

# Matrix algebra

Regular structure of MLPs facilitates algebraic shortcuts both in forward and backward pass.

$$f^k = W^k g^{k-1}.$$

Here  $f^k$  and  $g^{k-1}$  could be vectors if single datapoint, or matrices.

$$\frac{\partial L}{\partial g^{k-1}} = W^T \frac{\partial L}{\partial f^k}.$$

Again,  $\frac{\partial L}{\partial g^{k-1}}$  and  $\frac{\partial L}{\partial f^k}$  can be vectors or matrices.

# Matrix algebra

**Exercise:** Relate  $\frac{\partial L}{\partial f^k}$  and  $\frac{\partial L}{\partial g^k}$ , if  $g^k = \sigma(f^k)$  (coordinate-wise).

# Matrix algebra

**Exercise:** Relate  $\frac{\partial L}{\partial f^k}$  and  $\frac{\partial L}{\partial g^k}$ , if  $g^k = \sigma(f^k)$  (coordinate-wise).

**Solution:**

$$\frac{\partial L}{\partial f^k} = \frac{\partial L}{\partial g^k} \circ g^k \circ (1 - g^k).$$

$A \circ B$  is the Hadamard (element-wise) product.

# Matrix algebra

**Exercise:** Relate  $\frac{\partial L}{\partial f^k}$  and  $\frac{\partial L}{\partial W^k}$  (this is the matrix of partial derivatives over all elements of  $W^k$ ).

# Matrix algebra

**Exercise:** Relate  $\frac{\partial L}{\partial f^k}$  and  $\frac{\partial L}{\partial W^k}$  (this is the matrix of partial derivatives over all elements of  $W^k$ ).

**Solution:**

$$\frac{\partial L}{\partial W^k} = \frac{\partial L}{\partial f^k} \left( g^{k-1} \right)^T.$$

Note that for batches summing up happens automagically!

# The art of neural networks

In the non-modern setting, setting the parameters of neural networks was more art than science:

- ▶ Choosing the architecture.
- ▶ Choosing the units.
- ▶ Controlling the learning rate.
- ▶ Initializing the weights.
- ▶ And other (regularization, mini-batch size, etc.)


These issues were addressed by recent developments.

# Architecture


- ▶ Bigger network is more expressive.
- ▶ Bigger networks take longer to train.
- ▶ Deeper networks are much harder to train.
- ▶ Modern variants of SGD are much better at training deep nets. And we have stronger hardware.
- ▶ Confront: *If your deep neural net is not overfitting you should be using a bigger one!* – G. Hinton




# Units




sigmoid:  $\sigma(x) = \frac{1}{1+e^{-x}}$



tanh:  $\tanh(x) = \frac{e^x - e^{-x}}{e^x + e^{-x}}$



ReLU:  $\text{ReLU}(x) = \max(0, x)$



leaky ReLU:  $\text{LReLU}(x) = \max(ax, x)$

# Learning rate

- ▶ If the learning rate is too small, convergence is very slow.
- ▶ Learning rate that is too high results in oscillations.
- ▶ Some modern variants of SGD can automatically adapt.

# Initialization

- ▶ Very small initial weights slow down learning initially, or even completely vanish.
- ▶ Very large initial weights can lead to oscillations and random-like behaviour or even blow-up with no regularization.
- ▶ Heuristics, e.g. Glorot initialization: draw a weight from normal distribution with variance  $\frac{2}{n_{in} + n_{out}}$ .
- ▶ Modern techniques like batch normalization make initialization less of an issue.

# Recap

- ▶ Perceptron and its biological and computational motivation.
- ▶ Multilayer perceptron and how to compute its gradients fast.
- ▶ Tricky decisions when tuning MLPs.

# Recap

- ▶ Perceptron and its biological and computational motivation.
- ▶ Multilayer perceptron and how to compute its gradients fast.
- ▶ Tricky decisions when tuning MLPs.

But do not worry:

- ▶ Sometimes MLPs work well even without excessive tuning.
- ▶ Next lecture, we will make things much more stable.